



DISTRIBUTED FILE SYSTEM

DILIP KUMAR SHRESTHA

DILIP KUMAR SHRESTHA: GCES

INTRODUCTION

A distributed file system enables programs to store and access remote files exactly as they do local ones, allowing users to access files from any computer on a network. The performance and reliability experienced for access to files stored at a server should be comparable to that for files stored on local disks.

STORAGE SYSTEMS AND THEIR PROPERTIES

	<i>Sharing</i>	<i>Persistence</i>	<i>Distributed cache/replicas</i>	<i>Consistency maintenance</i>	<i>Example</i>
Main memory	×	×	×	1	RAM
File system	×	✓	×	1	UNIX file system
Distributed file system	✓	✓	✓	✓	Sun NFS
Web	✓	✓	✓	×	Web server
Distributed shared memory	✓	×	✓	✓	Ivy (DSM, Ch. 6)
Remote objects (RMI/ORB)	✓	×	×	1	CORBA
Persistent object store	✓	✓	×	1	CORBA Persistent State Service
Peer-to-peer storage system	✓	✓	✓	2	OceanStore (Ch. 10)

DILIP KUMAR SHRESTHA: GCES

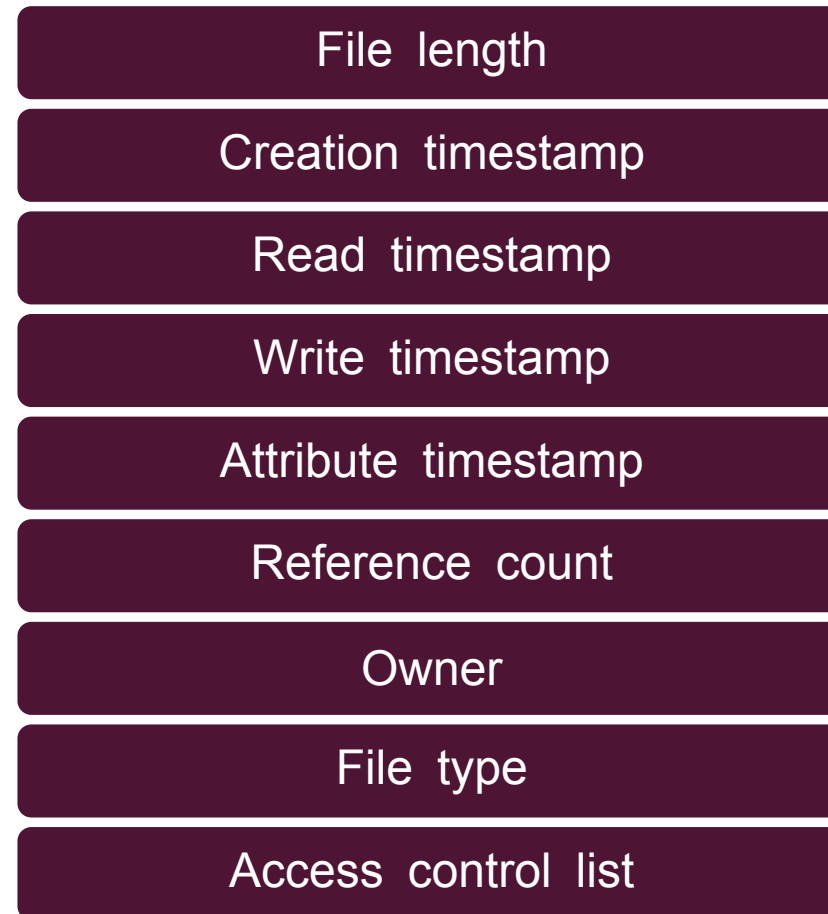
Types of consistency:

1: strict one-copy ✓: slightly weaker guarantees 2: considerably weaker guarantees

FILE SYSTEM MODULES

Directory module:	relates file names to file IDs
File module:	relates file IDs to particular files
Access control module:	checks permission for operation requested
File access module:	reads or writes file data or attributes
Block module:	accesses and allocates disk blocks
Device module:	performs disk I/O and buffering

FILE ATTRIBUTE RECORD STRUCTURE



DISTRIBUTED FILE SYSTEM REQUIREMENTS

Transparency

- Access transparency
- Location transparency
- Mobility transparency
- Performance transparency
- Scaling transparency

Concurrent file updates

File replication

Heterogeneity

Fault tolerance

Consistency

Security

Efficiency

FILE SERVICE ARCHITECTURE

An architecture that offers a clear separation of the main concerns in providing access to files is obtained by structuring the file service as three components – a flat file service, a directory service and a client module.

Flat file service: concerned with implementing operations on the contents of files (works on UFIDs)

Directory service: provides a mapping between text names for files and their UFIDs

Client module: runs in each client computer, integrating and extending the operations of the flat file service and the directory service

FILE OPERATIONS: UNIX

filedes = *open(name, mode)*

filedes = *creat(name, mode)*

Opens an existing file with the given *name*.

Creates a new file with the given *name*.

Both operations deliver a file descriptor referencing the open file. The *mode* is *read*, *write* or both.

status = *close(filedes)*

Closes the open file *filedes*.

count = *read(filedes, buffer, n)*

Transfers *n* bytes from the file referenced by *filedes* to *buffer*.

count = *write(filedes, buffer, n)*

Transfers *n* bytes to the file referenced by *filedes* from *buffer*.

Both operations deliver the number of bytes actually transferred and advance the read-write pointer.

pos = *lseek(filedes, offset, whence)*

Moves the read-write pointer to *offset* (relative or absolute, depending on *whence*).

status = *unlink(name)*

Removes the file *name* from the directory structure. If the file has no other names, it is deleted.

status = *link(name1, name2)*

Adds a new name (*name2*) for a file (*name1*).

status = *stat(name, buffer)*

Puts the file attributes for file *name* into *buffer*.

FILE OPERATIONS: FLAT FILE SERVICE

Read(FileId, i, n) → Data
— throws *BadPosition*

If $1 \leq i \leq \text{Length}(\text{File})$: Reads a sequence of up to n items from a file starting at item i and returns it in *Data*.

Write(FileId, i, Data)
— throws *BadPosition*

If $1 \leq i \leq \text{Length}(\text{File})+1$: Writes a sequence of *Data* to a file, starting at item i , extending the file if necessary.

Create() → FileId

Creates a new file of length 0 and delivers a UFID for it.

Delete(FileId)

Removes the file from the file store.

GetAttributes(FileId) → Attr

Returns the file attributes for the file.

SetAttributes(FileId, Attr)

Sets the file attributes (only those attributes that are not shaded in Figure 12.3).

FILE OPERATIONS: FLAT DIRECTORY SERVICE

Lookup(Dir, Name) → FileId
— throws *NotFound*

Locates the text name in the directory and returns the relevant UFID. If *Name* is not in the directory, throws an exception.

AddName(Dir, Name, FileId)
— throws *NameDuplicate*

If *Name* is not in the directory, adds (*Name, File*) to the directory and updates the file's attribute record.
If *Name* is already in the directory, throws an exception.

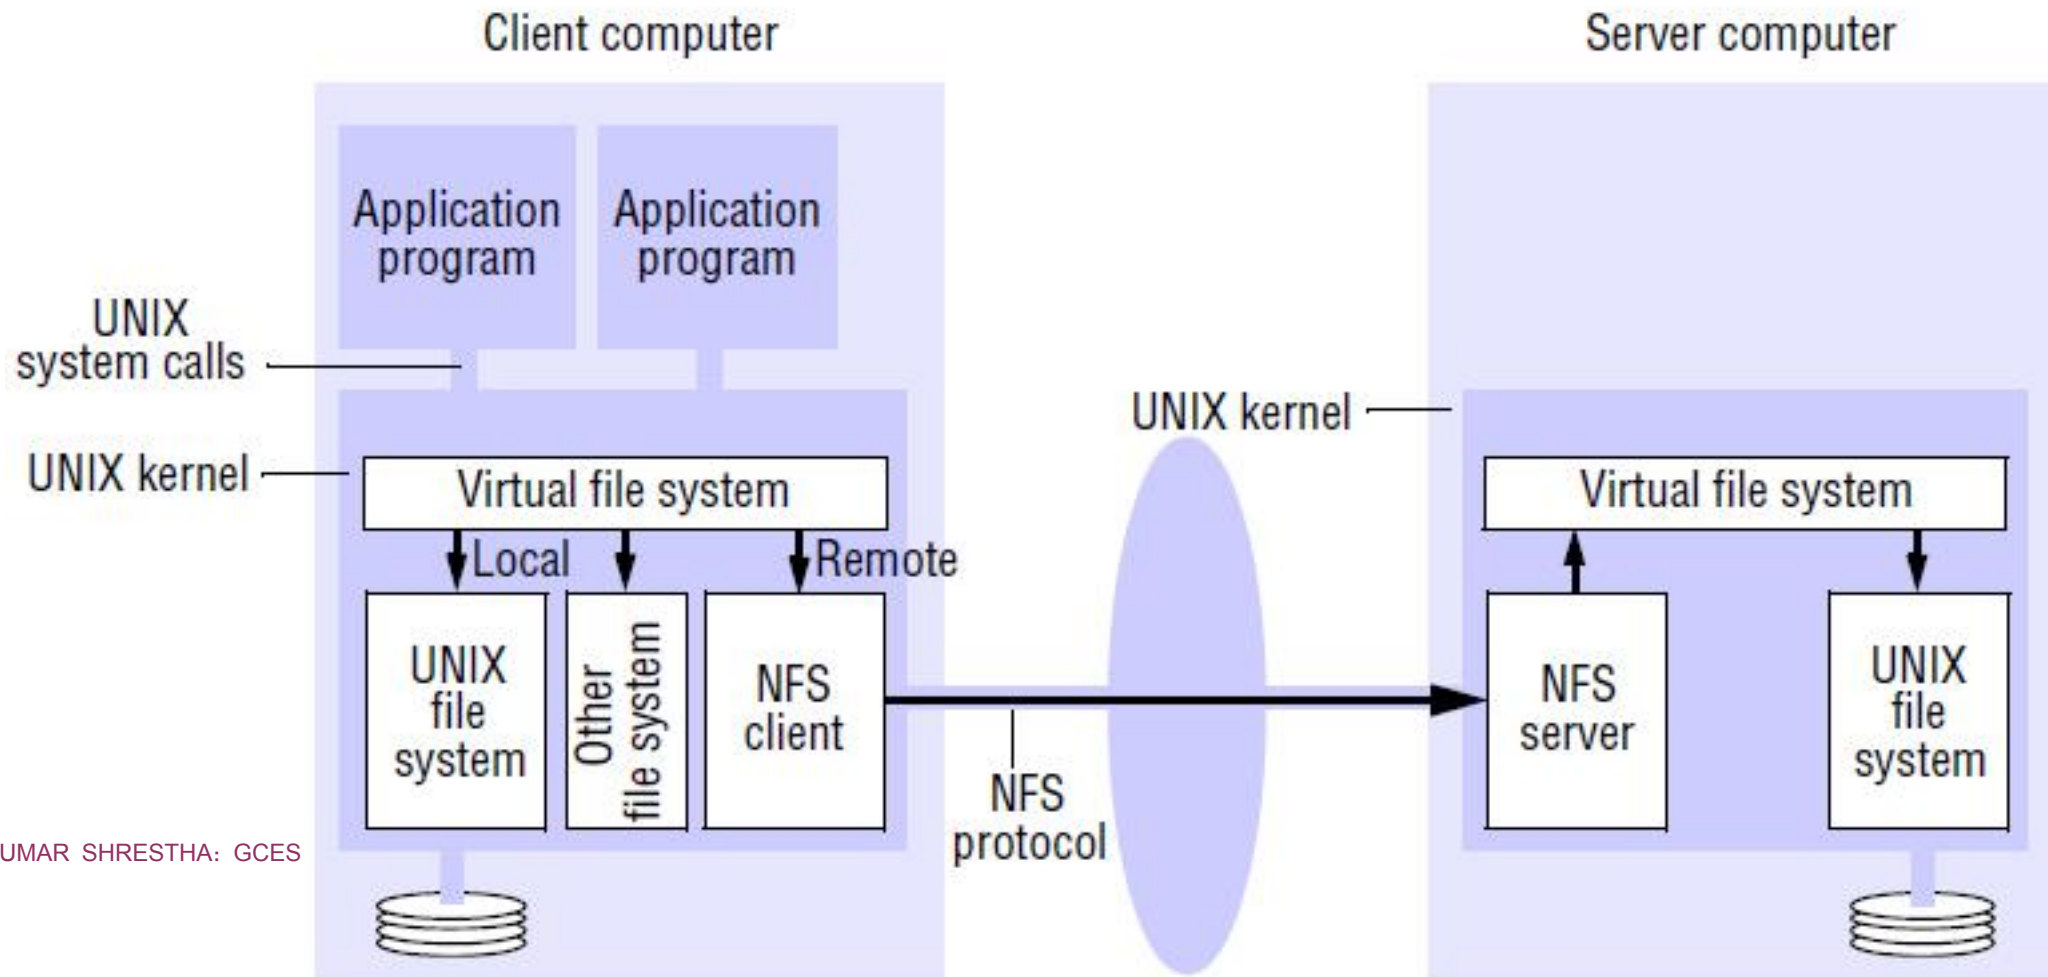
UnName(Dir, Name)
— throws *NotFound*


If *Name* is in the directory, removes the entry containing *Name* from the directory.
If *Name* is not in the directory, throws an exception.

GetNames(Dir, Pattern) → NameSeq

Returns all the text names in the directory that match the regular expression *Pattern*.

SUN NETWORK FILE SYSTEM





DILIP KUMAR SHRESTHA: GCES

GENERAL USAGES

TERMINOLOGIE S

Virtual File System

Client Integration

Access control and authentication

NFS server interface

Mount service

Path name translation

Automounter

Server caching

Client caching

SERVER OPERATIONS

- GETATTR: Get file attributes
- SETATTR: Set file attributes
- LOOKUP: Lookup filename
- ACCESS: Check access permission
- READLINK: Read from symbolic link
- READ: Read from file
- WRITE: Write to file
- CREATE: Create a file
- MKDIR: Create a directory
- SYMLINK: Create a symbolic link
- MKNOD: Create a special device
- REMOVE: Remove a file
- RMDIR: Remove a directory
- RENAME: Rename a file or directory
- LINK: Create link to an object
- REaddir: Read From directory
- REaddirplus: Extended read from directory
- FSSTAT: Get dynamic file system information
- FSINFO: Get static file system information
- PATHCONF: Retrieve POSIX information

USEFUL REFERENCES

Protocol v3 specification

- <https://tools.ietf.org/html/rfc1813>

Paper

- <http://people.cs.pitt.edu/~manas/courses/2510/nfs3.pdf>

Installation

- <https://help.ubuntu.com/lts/serverguide/network-file-system.html.en>

THANKYOU

